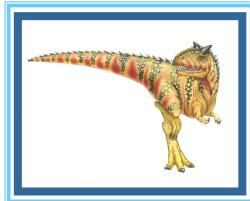


I. Operating Systems: Structure and Principles

Customized by Dr. Zhu for CS3600 at MSU Denver



Operating System Concepts – 10th Edition

Silberschatz, Galvin and Gagne ©2018

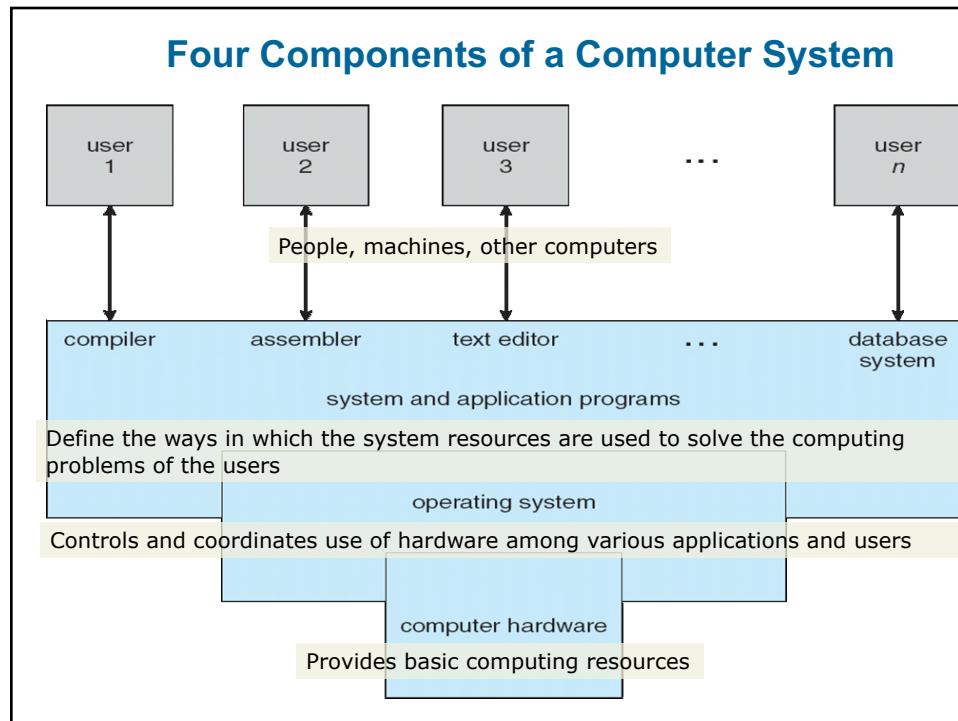
1

Topics

Reading Materials: Chapters 1 and 2 in the textbook

- What is an operating system
- Services/functionality of an operating system
- OS modes: batching, multitasking, and time-sharing
- Influences of security, networking, multimedia, windows
- Abstractions, virtual machines
- Structure methods (monolithic, layered, modular, micro-kernel, objective oriented modes)

2



3

Operating System Definition

- OS is a **resource allocator**
 - Manages all resources
 - Decides between conflicting requests for efficient and fair resource use
- OS is a **control program**
 - Controls execution of programs to prevent errors and improper use of the computer

4

Operating System Definition (Cont.)

- No universally accepted definition
- “Everything a vendor ships when you order an operating system” is a good approximation
 - But varies wildly
- “The one program running at all times on the computer” is the **kernel**.
- Everything else is either
 - a **system program** (ships with the operating system) , or
 - an **application program**.

5

Operating System Services/Functionality

- Operating systems provide an environment for execution of programs and services to programs and users
- One set of operating-system services provides functions that are helpful to the **user**:
 - **User interface** - Almost all operating systems have a user interface (**UI**).
 - ▶ Varies between **Batch** (early computers were batch systems processing jobs in bulk w/ predetermined input from files or other data sources), **Command-Line (CLI)**, text commands and a method for entering text), **Graphics User Interface (GUI**, a window system w/ a mouse), **Touch-Screen Interface** (allow user to slide fingers across the screen or press buttons on the screen)
 - **Program execution** - The system must be able to **load** a program into memory, to **run** that program, and to **end** execution, either normally or abnormally (indicating error)
 - **I/O operations** - A running program may require I/O, which may involve a file or an I/O device while users usually cannot directly control I/O devices.

6

Operating System Services/Functionality (Cont.)

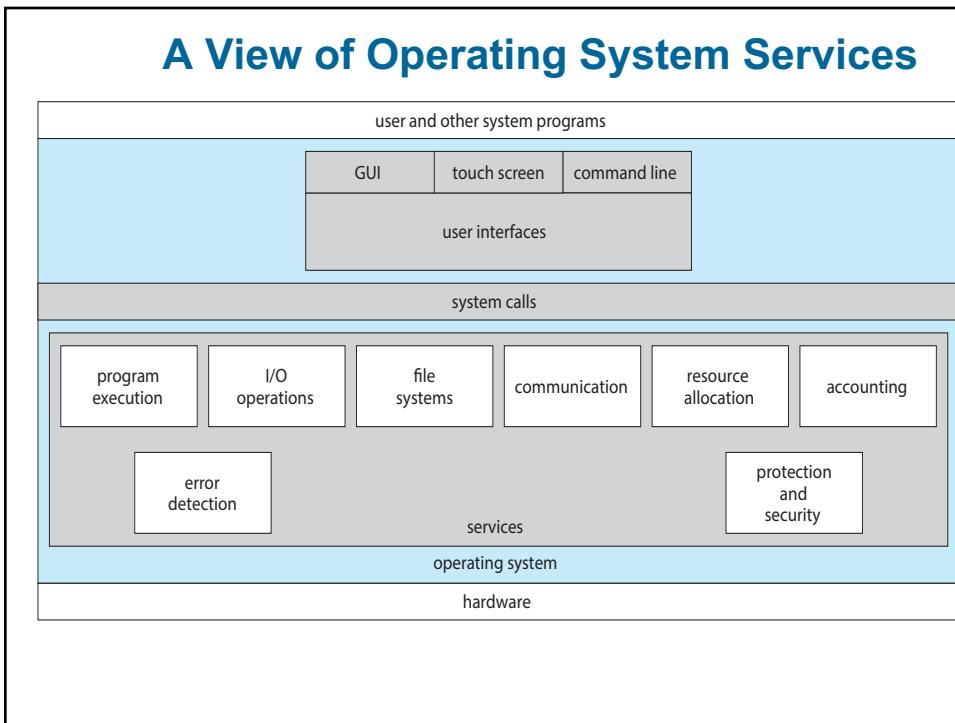
- One set of operating-system services provides functions that are helpful to the **user** (Cont.):
 - **File-system manipulation** - The file system is of particular interest. Programs need to read and write files and directories, create and delete them, search them, list file information, permission management.
 - **Communications** – Processes may exchange information, on the same computer or between computers over a network
 - ▶ Communications may be via **shared memory** or through **message passing** (packets moved between processors or between computers)
 - **Error detection** – OS needs to be constantly aware of possible errors
 - ▶ May occur in the CPU and memory hardware, in I/O devices, in user program
 - ▶ For each type of error, OS should take the appropriate action (e.g., to terminate an error-causing process, to return an error code, or to halt the system) to ensure correct and consistent computing
 - ▶ Debugging facilities can greatly enhance the user's and programmer's abilities to efficiently use the system

7

Operating System Services/Functionality (Cont.)

- Another set of OS functions exists for ensuring the efficient operation of the **system itself** via resource sharing
 - **Resource allocation** - When multiple users or multiple jobs running concurrently, resources must be allocated to each of them
 - ▶ Many types of resources - CPU cycles, main memory, file storage, I/O devices
 - **Accounting** - To keep track of which users use how much and what kinds of computer resources
 - **Protection and security** - The owners of information stored in a multiuser or networked computer system may want to control use of that information, concurrent processes should not interfere with each other
 - ▶ **Protection** involves ensuring that all access to system resources is controlled
 - ▶ **Security** of the system from outsiders requires **user authentication**, extends to defending external I/O devices from invalid access attempts

8



9

System Calls

- Demo on cs3600a.msudenver.edu
 - Enter a directory of your choice and assume there is a file named `srcFile.txt`
 - Run the following command to display all the system calls to copy a source file to a destination file


```
strace cp srcFile.txt dstFile.txt
```
 - Linux systems provide the `strace` utility to trace system calls
 - Read Section 2.3 for more details on “System Calls”

10

Computer and Operating System Startup

- **Bootstrap program** is loaded at power-up or reboot
 - Typically stored in ROM or EPROM, generally known as **firmware**
 - Initializes all aspects of system
 - Loads **operating system kernel** and starts execution
- **The kernel** starts providing services to the system and its users
 - Some services are provided outside of the kernel by system programs that are loaded into memory at boot time to become **system daemons**
- **System daemons** run the entire time the kernel is running
 - On Linux, the first system program is “**systemd**”, which starts other daemons
- Once this phase (loading/executing system daemons) is complete, the system is fully booted and waits for **some event** to occur.

11

Operating-System Operations

- **Events** are almost always signaled by the occurrence of an **interrupt**
- **Interrupt driven** (hardware and software)
 - **Hardware interrupt** by one of the devices
 - **Exception** or **Trap**: A **software-generated interrupt** caused by
 - ▶ A software error (e.g., division by zero, invalid memory access)
 - ▶ A specific request from a user program, which executes a special operation called **system call**, for an operating-system service
 - ▶ Or other process problems including an infinite loop, processes modifying each other or the operating system

12

Operating-System Operations (Cont.)

- **Multiprogramming (Batch system)** needed for efficiency
 - A single user/program cannot keep CPU and I/O devices busy at all times
 - Users typically *want* to run more than one program at a time
 - Multiprogramming organizes jobs (code and data) so CPU always has one to execute
 - A subset of total jobs in system is kept in memory
 - One job selected and run via **job scheduling**
 - When it has to wait (for I/O for example), OS switches to another job
- **Timesharing (multitasking)** is logical extension in which CPU switches jobs so frequently that users can interact with each job while it is running, creating **interactive** computing
 - **Response time** should be < 1 second
 - Each user has at least one program executing in memory ⇒ **process**
 - If several jobs ready to run at the same time ⇒ **CPU scheduling**
 - If processes don't fit in memory, **swapping** moves them in and out to run
 - **Virtual memory** allows execution of processes not completely fit in memory

13

Memory Layout for Multiprogrammed System



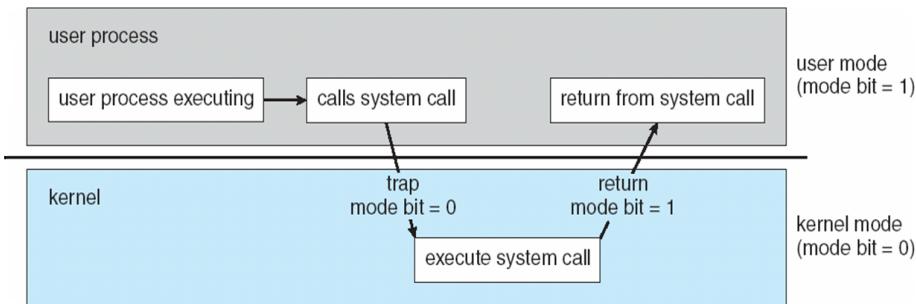
14

Operating-System Modes

- **Dual-mode** operation allows OS to protect itself and other system components
 - **User mode** and **kernel mode**
 - **Mode bit** provided by hardware
 - ▶ Provides ability to distinguish when system is running user code or kernel code
 - ▶ Some instructions designated as **privileged**, only executable in **kernel mode**
 - ▶ **System call** changes mode to **kernel**, return from call resets it to **user**
- Increasingly CPUs support **multi-mode** operations, e.g.,
 - ARM general processors except in the M-profile: ARMv4, ARMv7-A, etc.
 - **virtual machine manager (VMM)** mode for guest **VMs**

15

Transition from User to Kernel Mode



- **Timer** to prevent infinite loop / process hogging resources
 - Timer is set to interrupt the computer after some time period
 - Keep a counter that is decremented by the **physical clock**.
 - Operating system set the counter (privileged instruction)
 - When counter zero generate an **interrupt**
 - Set up before scheduling process to regain control or terminate program that exceeds allotted time

16

Protection and Security

- **Protection** – any mechanism for controlling access of processes or users to resources defined by the OS
- **Security** – defense of the system against internal and external attacks
 - Huge range, including denial-of-service, worms, viruses, identity theft, theft of service
- Systems generally first distinguish among users, to determine who can do what
 - User identities (**user IDs**, security IDs) include name and associated number, one per user
 - User ID then associated with *all files, processes of that user* to determine **access control**
 - Group identifier (**group ID**) allows set of users to be defined and controls managed, then also associated with **each process, file**
 - **Privilege escalation** allows user to change to effective ID with more rights
 - ▶ On UNIX, for instance, the **setuid** attribute on a **program** causes that program to run with **the user ID of the owner of the file**, rather than the current user's ID

17

Computing Environments - Traditional

- Stand-alone general purpose machines
- But blurred as most systems **interconnect** with others (i.e., **the Internet**)
- **Portals** provide web access to internal systems
 - Gateways between requestors and services running on provider computers
- **Network computers (thin clients)** are like Web terminals
 - Limited computers that only understand web-based computing
- Mobile computers interconnect via **wireless networks**
- Networking becoming ubiquitous – even home systems use **firewalls** to protect home computers from Internet attacks

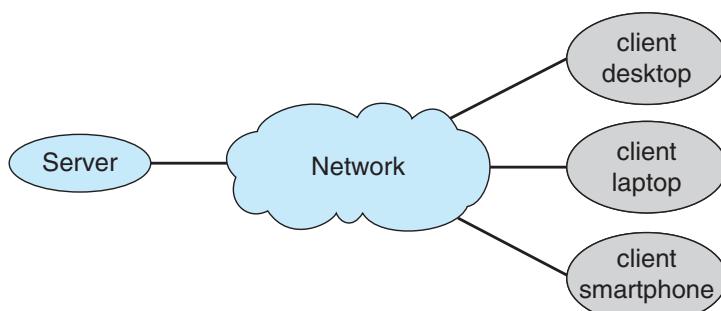
18

Computing Environments - Mobile

- Handheld smartphones, tablets, etc.
- The functional difference between them and a “traditional” laptop?
 - Features on mobile devices have become so rich that the distinction in functionality between may be difficult to discern
 - In fact, a contemporary mobile device can now provide functionality that is either unavailable or impractical on a laptop computer
- Extra feature – more OS features (GPS, gyroscope)
- Allows new types of apps like ***augmented reality***
- Use IEEE 802.11 wireless, or cellular data networks for connectivity
- Leaders are [Apple iOS](#) and [Google Android](#)

19

Computing Environments – Client-Server

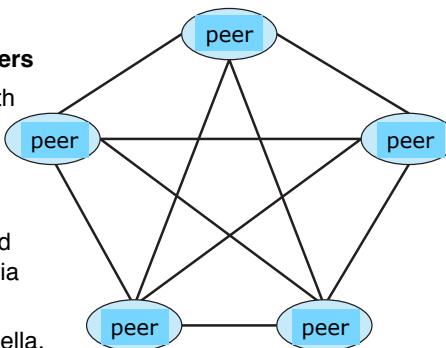


- Client-Server Computing: a form of specialized ***distributed computing***
 - Dumb terminals supplanted by smart PCs
 - Many systems now **servers**, responding to requests generated by **clients**
 - ▶ **Compute-server system** provides an interface to client to request services (i.e., database)
 - ▶ **File-server system** provides interface for clients to store and retrieve files

20

Computing Environments - Peer-to-Peer

- Another model of **distributed system**
- P2P does not distinguish clients and servers
 - Instead all nodes are considered **peers**
 - May each act as client, server or both
 - Node must join P2P network
 - ▶ Registers its service with central lookup service on network, or
 - ▶ Broadcast request for service and respond to requests for service via **discovery protocol**
 - Examples include Napster and Gnutella, **Voice over IP (VoIP)** such as Skype



21

Computing Environments - Virtualization

- Allows operating systems to run applications within other OSes
 - Vast and growing industry
- **Emulation** used when **source CPU type** different from **target type** (i.e. when Apple switched from PowerPC to Intel x86 for its desktops and laptops, "Rosetta" was included in the OS to allow applications compiled for PowerPC to run on Intel x86.)
 - Can be extended to allow an entire operating system written for one CPU to run on another
 - Generally slowest method (every instruction needs to be translated.)
 - When computer language not compiled to native code – **Interpretation**
- **Virtualization** – OS natively compiled for CPU, running **guest OSes** also natively compiled for the same CPU
 - **VMM** (virtual machine Manager) provides virtualization services

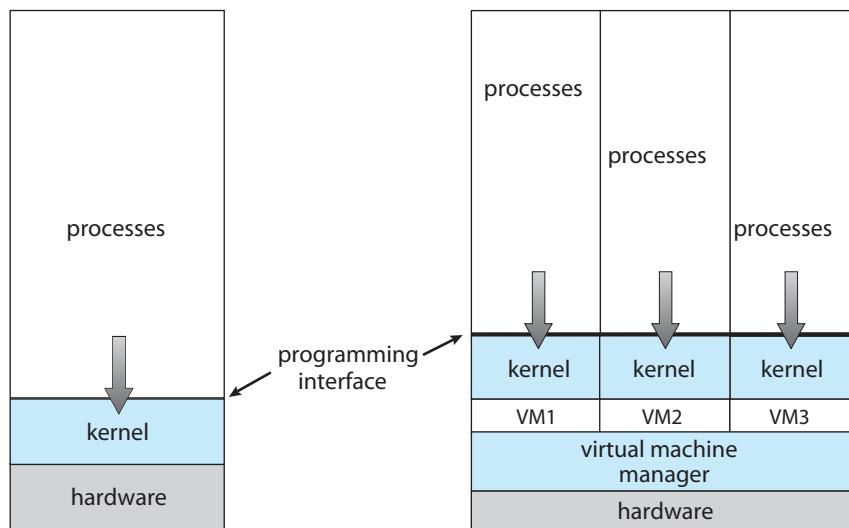
22

Virtualization (Cont.)

- Use cases involve laptops and desktops running multiple OSes for exploration or compatibility
 - Apple laptop running Mac OS X **host**, Windows 10 as a **guest**
 - Developing apps for multiple OSes without having multiple systems
 - QA testing applications without having multiple systems
 - Executing and managing compute environments within data centers
- Some VMMs like *VMware ESX* and *Citrix XenServer* no longer run on host operating systems but rather are the host operating systems
 - Provide services and resource management to virtual machine processes
 - There is no general purpose host then

23

Virtualization w/o General Purpose Host OS



24

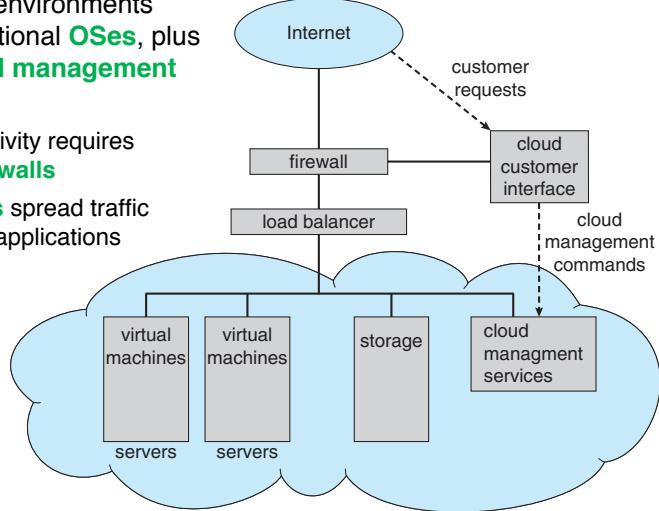
Computing Environments – Cloud Computing

- Delivers computing, storage, even apps as a service across a network
- Logical extension of virtualization because it uses virtualization as the base for its functionality.
 - Amazon **EC2** has thousands of servers, millions of virtual machines, petabytes of storage available across the Internet, pay based on usage
- Many types
 - **Public cloud** – available via Internet to anyone willing to pay
 - **Private cloud** – run by a company for the company's own use
 - **Hybrid cloud** – includes both public and private cloud components
 - Software as a Service (**SaaS**) – one or more applications available via the Internet (i.e., word processor)
 - Platform as a Service (**PaaS**) – software stack ready for application use via the Internet (i.e., a database server)
 - Infrastructure as a Service (**IaaS**) – servers or storage available over Internet (i.e., storage available for backup use)

25

Computing Environments – Cloud Computing (Cont.)

- Cloud computing environments composed of traditional **OSes**, plus **VMs**, plus **cloud management tools**
 - Internet connectivity requires security like **firewalls**
 - **Load balancers** spread traffic across multiple applications



26

Computing Environments – Real-Time Embedded Systems

- Real-time embedded systems: most prevalent form of computers
 - Vary considerable, special purpose, limited purpose OS, **real-time OS**
 - Use expanding
- Real-time OS has well-defined fixed time constraints
 - Processing **must** be done within constraint
 - Correct operation only if constraints met
- Many other special computing environments as well
 - Some have OSes, some perform tasks without an OS

27

Operating System Structure

- General-purpose OS is very large program
- Various ways to structure ones
 - Monolithic – the original UNIX
 - Layered – an abstraction
 - Microkernel – Mach
 - Modules – Modern UNIX (Linux, macOS, and Solaris)

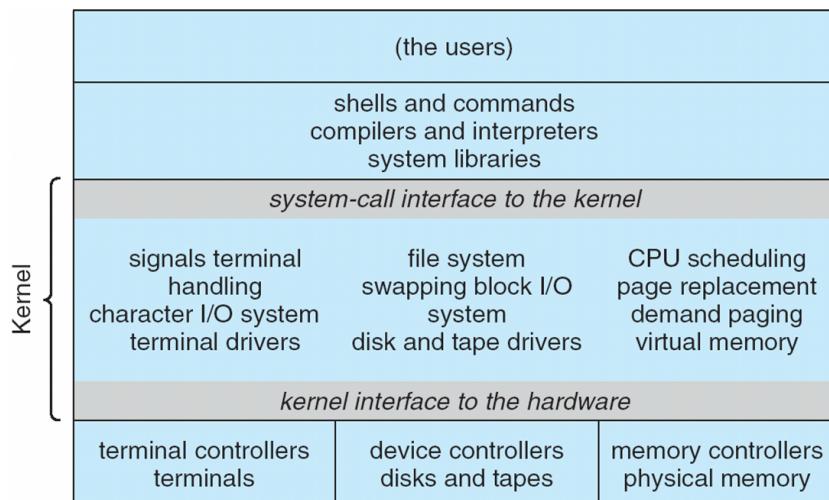
28

Monolithic Structure -- UNIX

- **Monolithic structure:** place all of the functionality of the kernel into a single, static binary file that runs in a single address space.
 - tightly **coupled system:** changes to one part of the system can have wide-ranging effects on other parts
 - a **distinct performance advantage:** very little overhead in the system-call interface, and fast communication within the kernel
- **UNIX** – limited by hardware functionality, the original UNIX operating system has such limited structure and consists of two separable parts
 - **System programs**
 - **The kernel**
 - ▶ Consists of everything *below* the system-call interface and *above* the physical hardware
 - ▶ Provides the *file system, CPU scheduling, memory management, and other operating-system functions*
 - ▶ Combine an enormous amount of functionality into *one single address space in one level*

29

Traditional UNIX System Structure

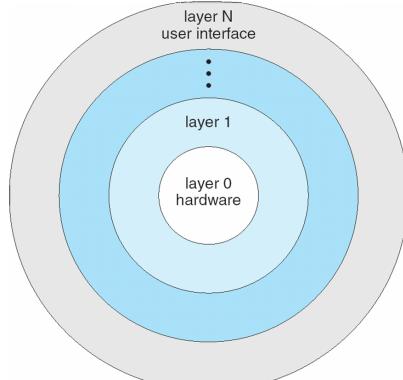


- Beyond simple but not fully layered
- As UNIX expanded, the kernel becomes large and difficult to manage

30

Layered Approach for Modularity

- The operating system is divided into a number of layers, each built on top of lower layers.
 - The bottom layer (layer 0), is the hardware; the highest (layer N) is the user interface.
 - Each layer consists of data structures and a set of functions that can be invoked by higher layers
 - Each layer can invoke operations on lower layers.
 - Main advantage: simplicity of construction and debugging
- Layered systems have successfully used in computer networks (such as TCP/IP) and web applications. Few operating systems use a pure layered approach.
 - Challenging to appropriately define the functionality of each layer
 - Poor overall performance due to the overhead of requiring a user program to traverse through multiple layers to obtain an OS service



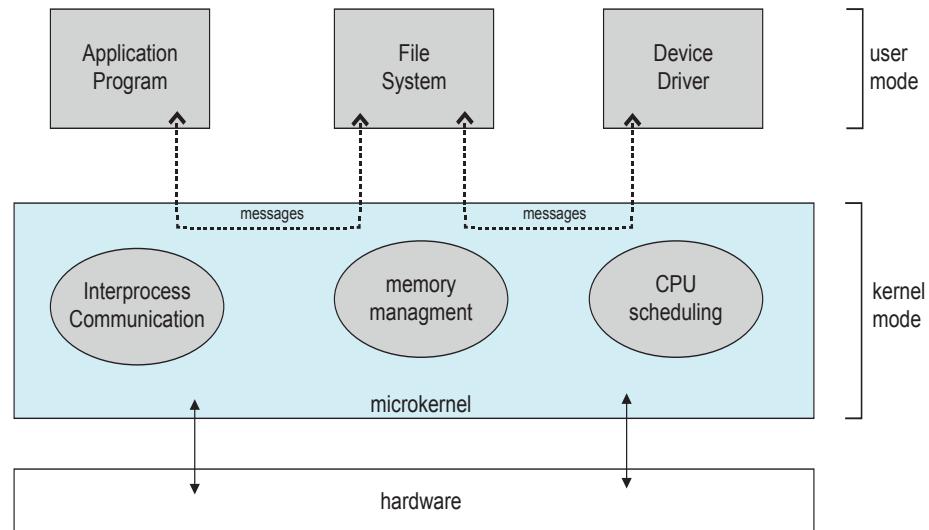
31

Microkernel System Structure

- Remove all nonessential components from the kernel and implement them as user-level programs that reside in separate address spaces
- **Mach**: proposed in the mid-1980s by researchers to modularize the kernel using the **microkernel** approach
 - Mac OS X kernel (**Darwin**) for **macOS** and **iOS**: partly based on **Mach**
 - ▶ the best-known illustration of microkernel
- User modules never directly communicate, but indirectly by exchanging messages with the microkernel via **message passing**
- Benefits:
 - Easier to extend a microkernel
 - Easier to port the operating system to new architectures
 - More reliable (less code is running in kernel mode)
 - More secure
- Detriments:
 - Performance overhead of user space to kernel space communication

32

Microkernel System Structure (Cont.)



33

Object-Oriented Modules

- Many modern operating systems implement **loadable kernel modules**
 - The kernel has a set of core components and can link in additional services via modules, at boot time or during run time.
 - An object-oriented approach
 - Each core component is separate
 - Each talks to the others over **known interfaces**
 - Each is loadable **as needed** within the kernel
- Like a **layered** system, each module has defined, protected interfaces
 - but more flexible, because any module can call any other module.
- Like the **microkernel** approach, the primary module has only core functions and knowledge of load and communicate with other modules
 - but more efficient, because modules do not need to invoke message passing via the kernel in order to communicate.
 - Modern UNIX such as Linux, Solaris, and macOS, as well as Windows

34

Kernel Modules

■ Demo on cs3600a.msudenver.edu

- Run the following command to list all kernel modules that are currently loaded

```
lsmod
```

```
lsmod | tee curKernelModules
```

- ▶ “Size”: the size of the module (not memory used)
- ▶ “Used by”: the number of times the module is currently in use by running programs
 - next to this number is a list of other modules which refer to this one

- Read the programming project “Introduction to Linux Kernel Modules” at the end of Chapter 2 for more details